arXiv:2610.40745v1 [cs.DL] 29 Oct 2026

What It Costs to Trust a Source: A Reproducible Method for Institutional Assessment of Author Disambiguation

Simon J. Porter [1, *]

[1] *Digital Science, 6 Briset Street, London EC1M 5NR, United Kingdom*

(Dated: 29 October 2026)

Records of research outputs are curated with care by universities, and the effort is shaped by the external sources from which candidate records are drawn. The curation required before the author identities in such sources can be trusted is examined, and a reproducible method for estimating its cost is set out. Four tests are described; two can be run on public data alone, and two are conducted against an institution's own curated record, taken from Symplectic Elements Discovery portals. The method is demonstrated on OpenAlex and Dimensions across 30 institutions. Of the publications attributed to matched researchers, 50.4% were confirmed by the institutional record for OpenAlex and 69.6% for Dimensions, and fewer trusted researcher profiles were retained from OpenAlex at a matched confidence threshold in each of the 30 institutions. Carried through to effort, integrating OpenAlex requires 1.5 times the person-days of review and entry of Dimensions even when each feed is filtered at the threshold that minimises the amount of review required.

Keywords: author disambiguation; research information management; CRIS; ORCID; OpenAlex; Dimensions

I. INTRODUCTION

Since the early 2000s, bibliographic sources on the web have helped universities to assemble records of the research of their staff. In most research intensive institutions Current Research Information Systems (CRIS) now support compliance and assessment submissions, open-access monitoring, annual reporting and the public profiles through which an institution presents its people to collaborators, funders and the press [7]. The outcomes based on these uses are consequential for both individuals and institutions and hence records are curated in collaboration and with care by research offices, librarians and researchers. This activity is both ongoing and significant and hence the ongoing cost of curation is a recurring concern in the literature on research information management [2].

The scholarly records upon which this curation is based have been built in layers: Crossref has registered DOIs and publication metadata that has been deposited by source publishers since 2000 [21] and has made efforts to include historic content; and DataCite, founded in 2009 as a separate DOI registration agency, brought datasets, and later other research outputs, into the DOI system [5]. ORCID launched in 2012 to give research participants a public record of their outputs that they maintain themselves [20]. In each of these registries an assertion is made by a party with direct knowledge of a fact concerning a scholarly record, and each claim carries clear provenance, in the sense of FAIR principle R1.2 [69]. To be explicit, we interpret this to mean not only that every metadata item should conform to the FAIR principles, but that within reason, the provenance of the assertion should be implicitly clear. That is to say that, for example, metadata vested on Crossref can be assumed to be findable by virtue of Crossref's data services; accessible through Crossref's use of standardised, open data formats; interoperable as a result of the aforementioned formats; and reusable not only due to Crossref's licence terms but also implicitly because the record is trustworthy, being as provided to Crossref by its originator. This trust in: i) faithful reproduction of the record from a source; ii) the ability to assume such provenance for a whole record without the need to check individually; and iii) the knowledge that records have not been added or deleted and that the data provided by Crossref constitutes an unfiltered and unredacted set is critical to all the layers that then derive from it.

A second layer of sources, among them Web of Science, Scopus and Dimensions [67] and, more recently, OpenAlex [50], constructs knowledge that the registries do not hold: algorithmic identifiers for researchers whose names appear only as text, affiliations resolved from address strings, and classifications by field.

In this context a *secondary source* is an index that derives assertions about authorship algorithmically from metadata that it did not itself author, whether or not those assertions are later curated by hand. By a *person-level record* we mean the attribution of research outputs to a specific individual, expressed as an author identifier together with the outputs gathered under it, as distinct from the document-

* s.porter@digital-science.com; https://orcid.org/0000-0002-6151-8423

level record of what was published, where and with what citations, on which most comparisons of sources have concentrated. There are two kinds of failure that are common but which are not equally visible. A missing research output in a person's record is a gap that someone eventually notices, whereas a researcher profile that contains another person's work looks to every downstream system like a richer profile, even if the researcher appears to be publishing in a different domain with different collaborators, this is not unusual - many researchers change fields during their career—it might also mark a transistion to a different field of research. These "richer" profiles have an increasingly consequential impact due to the growing class of applications, from research-intelligence platforms that match researchers to policy-engagement opportunities [43] to integrity screens that flag anomalous publication behaviour [27, 47], and all these applications broadly take author identities as given and author disambiguation if not as a solved problem then a sufficiently solved problem.

An institution that brings such a source into its CRIS therefore takes on work that is not usually included in the calculation of the cost of ownership. We call it the *integration cost*, and measure it in effort: the curation that must be layered on top of a source before its claims about the institution's own people can be acted upon. Both the level of curation and where that curation lands is a matter of choice - informed by institutional and national cultural norms, financial considerations, trust in systems, trust in data, local ethical and regulatory practices, risk tolerance, and intended data use. Effort spent correcting a secondary source repairs one copy of the record inside another organisation's database, whereas confirmed person–publication links maintained in the institution's own CRIS fix the institutional record and, where connected through ORCID, supply an authority from which every downstream index can draw.

This paper sets out a reproducible method for estimating the integration cost of a scientometric source from data that an adopting institution already holds. Two of its four tests are *source-level*, using only a source's output and the ORCID public data file to assess how likely the source is to conflate distinct researchers; two are *institutional*, reading the source against an institution's own curated record as a proxy for a golden dataset. The method is demonstrated on OpenAlex [50] and Dimensions [22, 24], two indexes built on open persistent-identifier infrastructure and openly available metadata, across 30 institutions whose records are exposed through Symplectic Elements Discovery portals. The results are expressed in person-days, for administrators reviewing a feed and for researchers reviewing their own suggestions.

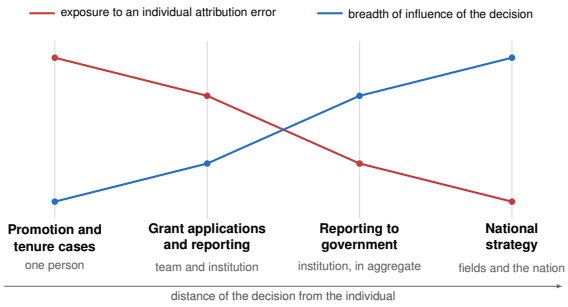


Figure 1: Uses of the institutional publication record, ordered by the distance of the decision from the individual researcher. The schematic is qualitative: the positions are ordinal and the crossing point is illustrative rather than measured.

The publication data held by a university are put to uses that differ considerably in how much error they can tolerate (Figure 1). At one extreme sits the case for promotion. In this case records are read by a committee whose decision bears directly on an individual's career; a misattributed or missing paper may materially affect the outcome and is a matter on which the individual might reasonably seek redress. The grant application is not far removed. A funder will expect an accurate account of the outcomes of earlier awards and of the work of each investigator named in the proposal, and an error reflects not only on the individuals concerned but on the standing of the institution with that funder.

Reporting to government occupies a middle position. Accuracy is still required, but most evaluative systems treat the data in aggregate, where an individual misattribution is diluted within the totals of a department or an institution; an exercise such as the UK's Research Excellence Framework, in which named outputs are submitted against named staff, sits rather closer to the stricter end than its aggregate presentation might suggest. At national level the data are used to discern trends: whether a field is growing, whether investment should follow, and how a research system aligns with a government's view of critical technologies. These are among the most consequential allocations made from such data, and yet the least sensitive to the accuracy of any individual record.

It follows that the bodies with the greatest influence over the direction of research funding are, in some sense, the least exposed to the consequences of inaccurate attribution. A national decision stands one step removed from the institutions that carry its consequences and two steps removed from the individuals whose records underlie it, while the cost of curation, and the risk that attends its absence, fall

on research offices and researchers.

Aggregation is forgiving of random error, though rather less so of systematic error, and profile conflation is of the latter kind. It concentrates on common names, which may be expected to cluster by naming convention and hence by country and region, and possibly by field, which are the units that national analyses compare, so whether a source's errors lean in one direction also matters beyond the research office. For the institution itself, the confidence required of a source becomes a property of the use to which the data are put, and a single institution will in practice need several levels of confidence from the same feed.

This paper is organised as follows: We provide a brief literature review (Section II) before describing the data and test methodology in Sec. III. Sec. IV contains our results, after which we discuss their implications in Sec. V) and note the limitations of our study in Sec. VI.

II. RELATED WORK

The cost of research information. The uses of a CRIS have expanded steadily, from compliance and assessment through open-access monitoring and research-data registration to public researcher profiles [7, 53]. Increasingly the CRIS also underpins what might be called *university-in-the-loop* reporting: measurement grounded in records that the institution itself verifies and warrants, in contrast to rankings and evaluations computed entirely from secondary sources outside its control. The Barcelona Declaration on Open Research Information reinforces this direction, committing its signatories to make openness the default for the research information they use, including to assess researchers, and to require that the data in research information systems can be exported and made open, so that the information used in evaluation is accessible and auditable by those being assessed [4], commitments that apply directly to institutional records such as those studied here.

A registry that drives public profiles and institutionally warranted reporting must be correct at the level of the individual researcher, and the literature is consistent that such correctness is bought with labour: profiling and cleansing routines [3], quality problems reported across national contexts [2], a standard of fitness for purpose that rises as uses accumulate [19], and work on integrating repositories and workflows [52]. What has been missing is a measurement connecting this literature to the choice of external source: how far the curation burden of a CRIS is determined in advance by the disambiguation quality of the source that feeds it, and whether that burden differs between sources consistently enough to be weighed in the choice.

The effort of record entry and verification. No published study, to our knowledge, measures how long it takes to accept or reject a publication that a CRIS suggests, so the effort of integration has to be assembled from adjacent tasks. For entering a record by hand, those who run a CRIS work with figures of a few minutes: the Portuguese national CV platform assumes three minutes per manually entered record in estimating the effort saved by its suggestion service [16], a case study of open-access workflows in Symplectic Elements estimates three minutes to create a record with its full text [58], and repository guidance allows five to ten minutes for a complete deposit [66].

Confirming an existing record has also been observed to be quicker than creating one. In a time-driven study of cataloguing, adapting an existing bibliographic description took 74 seconds against 306 seconds to create one, about three-quarters less, although across the whole workflow, physical processing included, the saving was a third [55]. Screening for systematic reviews supplies the nearest logged measurements of a quick bibliographic judgement: a mean of 35 seconds per abstract over nearly 400,000 decisions on one platform [26], and medians of 15 to 20 seconds in a small experiment reported by a vendor [35]. These are judgements of relevance against inclusion criteria rather than of authorship, and the full-text assessment that a doubtful case can require has been put, from one team's experience, at about fifteen minutes [65].

Administrative records may be maintained by researchers as well as administrators. A researcher who reviews suggestions personally also bears the overhead of logging into each session. Applying published second-factor times of roughly 10 to 21 seconds after the password to logs of more than 35 million two-factor logins at two universities, Reynolds et al. [51] find more than one login attempt in twenty abandoned or failed at the second step, and locate the burden less in the total time than in the position of the step on the critical path to the task the user came to do. Administration can also be considered to be an interruption to researchers' other responsibilities. Without supplying a cost per session, research on interruption finds that interrupted work is often not resumed for tens of minutes, during which other work intervenes [28, 32], and people compensate for interruption by working faster while reporting greater stress [33]. Estimating the effort that research-information workflows consume, and the share of it that persistent identifiers can remove, has also been developed at national scale [6].

Comparing scientometric sources. OpenAlex [50] inherited both the corpus and the ambition of